

# CLydonAssignment2

March 1, 2022

## 1 Assignment 1

- 1) Read chapters 2-4 in your textbook and review the coding examples we went over in class
- 2) Review the Keras documentation for things like the Layer types and Optimizer types to better familiarize yourself
- 3) Redo the coding examples and do the “Further experiments” in chapter 3 of your book with datasets you find interesting and may want to use for further assignments and projects
- 4) Provide a brief write up of the experiment and analysis.

## 2 table of contents

1. import libraries
2. data prep
  1. reading data
  2. finding summary data
  3. dummyizing data
  4. creating predictor & outcome attribute lists
  5. transforming data into numpy arrays
  6. creating training and testing data
  7. standardizing the data (z-scoring)
3. modeling
  1. initial model
    1. building layers
    2. fitting model on training data
    3. predicting with test data
  2. further experiments
    1. model using one hiddden layer
    2. model using more hidden units
    3. model using mae loss
    4. model using tanh activation
  3. kfold optimization
    1. build model function definition
    2. defining k-fold variables
    3. k-fold loop
    4. plotting k-fold results
    5. building model from validation

## 2.1 1 import libraries

```
[250]: # data prep
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

import numpy as np
import pandas as pd

# keras imports
from keras import models
from keras import layers
from keras import optimizers

# random seed for pseudo-random operations
random_seed = 100
```

## 2.2 2 data prep

### 2.1 reading dataset

```
[251]: math_data = pd.read_csv("https://raw.githubusercontent.com/connoralydon/
↳ datasets/main/student-math.csv", sep=";")
```

### 2.2 summary info about data

```
[252]: math_data.head
```

```
[252]: <bound method NDFrame.head of
Fedu      Mjob      Fjob      \
0         GP      F      18      U      GT3      A      4      4      at_home      teacher
1         GP      F      17      U      GT3      T      1      1      at_home      other
2         GP      F      15      U      LE3      T      1      1      at_home      other
3         GP      F      15      U      GT3      T      4      2      health      services
4         GP      F      16      U      GT3      T      3      3      other      other
..      ...      ..      ...      ...      ...      ...      ...      ...
390      MS      M      20      U      LE3      A      2      2      services      services
391      MS      M      17      U      LE3      T      3      1      services      services
392      MS      M      21      R      GT3      T      1      1      other      other
393      MS      M      18      R      LE3      T      3      2      services      other
394      MS      M      19      U      LE3      T      1      1      other      at_home

... famrel freetime goout Dalc Walc health absences G1 G2 G3
0      ...      4      3      4      1      1      3      6      5      6      6
1      ...      5      3      3      1      1      3      4      5      5      6
2      ...      4      3      2      2      3      3      10      7      8      10
3      ...      3      2      2      1      1      5      2      15      14      15
4      ...      4      3      2      1      2      5      4      6      10      10
..      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
```

|     |     |   |   |   |   |   |   |    |    |    |    |
|-----|-----|---|---|---|---|---|---|----|----|----|----|
| 390 | ... | 5 | 5 | 4 | 4 | 5 | 4 | 11 | 9  | 9  | 9  |
| 391 | ... | 2 | 4 | 5 | 3 | 4 | 2 | 3  | 14 | 16 | 16 |
| 392 | ... | 5 | 5 | 3 | 3 | 3 | 3 | 3  | 10 | 8  | 7  |
| 393 | ... | 4 | 4 | 1 | 3 | 4 | 5 | 0  | 11 | 12 | 10 |
| 394 | ... | 3 | 2 | 3 | 3 | 3 | 5 | 5  | 8  | 9  | 9  |

[395 rows x 33 columns]>

[253]: `math_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
#   Column          Non-Null Count  Dtype
---  -
0   school          395 non-null    object
1   sex              395 non-null    object
2   age              395 non-null    int64
3   address          395 non-null    object
4   famsize          395 non-null    object
5   Pstatus          395 non-null    object
6   Medu             395 non-null    int64
7   Fedu             395 non-null    int64
8   Mjob             395 non-null    object
9   Fjob             395 non-null    object
10  reason           395 non-null    object
11  guardian         395 non-null    object
12  traveltime       395 non-null    int64
13  studytime        395 non-null    int64
14  failures         395 non-null    int64
15  schoolsup        395 non-null    object
16  famsup           395 non-null    object
17  paid             395 non-null    object
18  activities       395 non-null    object
19  nursery          395 non-null    object
20  higher           395 non-null    object
21  internet         395 non-null    object
22  romantic         395 non-null    object
23  famrel           395 non-null    int64
24  freetime         395 non-null    int64
25  goout            395 non-null    int64
26  Dalc             395 non-null    int64
27  Walc             395 non-null    int64
28  health           395 non-null    int64
29  absences         395 non-null    int64
30  G1               395 non-null    int64
31  G2               395 non-null    int64
32  G3               395 non-null    int64
```

```
dtypes: int64(16), object(17)
memory usage: 102.0+ KB
```

```
[254]: math_data.dtypes
```

```
[254]: school      object
sex            object
age           int64
address       object
famsize       object
Pstatus       object
Medu          int64
Fedu          int64
Mjob          object
Fjob          object
reason        object
guardian      object
traveltime    int64
studytime     int64
failures      int64
schoolsup     object
famsup        object
paid          object
activities    object
nursery       object
higher        object
internet      object
romantic      object
famrel        int64
freetime      int64
goout         int64
Dalc          int64
Walc          int64
health        int64
absences      int64
G1            int64
G2            int64
G3            int64
dtype: object
```

```
[255]: math_data.columns
```

```
[255]: Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
        'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
        'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
        'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
        'Walc', 'health', 'absences', 'G1', 'G2', 'G3'],
        dtype='object')
```

```
[256]: math_data.shape
```

```
[256]: (395, 33)
```

```
[257]: # observe the different values used for species - 2 different values: 0, 1
def print_range(data):
    print(f"range of {data.min()} to {data.max()}\n")

print("First test")
print_range(math_data["G2"])

print("Second test")
print_range(math_data["G1"])

print("Final test")
print_range(math_data["G3"])
```

```
First test
range of 0 to 19
```

```
Second test
range of 3 to 19
```

```
Final test
range of 0 to 20
```

```
[258]: print("total null values")
math_data.isnull().sum(axis=0).sum()
```

```
total null values
```

```
[258]: 0
```

```
[259]: print("num duplicates")
math_data.duplicated().sum(axis=0)
```

```
num duplicates
```

```
[259]: 0
```

### *2.3 dummyizing data*

```
[260]: categorical_cols = {}

for col in math_data.columns:
    if(math_data[col].dtype == "object"):
        categorical_cols[col] = list(math_data[col].unique())
```

```
numerical_cols = list(set(math_data.columns) - set(categorical_cols.keys())) #  
↳ set difference  
  
math_data_dummy = pd.get_dummies(math_data, columns = categorical_cols.keys())
```

**2.4 model variables** removing g1, g2, g3 because they would be highly autocorrelary. someone who does good on one test is likely to do good on another.

```
[261]: predictors = list(math_data_dummy.columns)  
predictors.remove("G1") # first exam  
predictors.remove("G2") # second exam  
predictors.remove("G3") # third exam  
  
outcome = "G3"
```

**2.5 transforming data into numpy arrays** this is because keras won't take pandas dataframes

```
[262]: X = np.array(math_data_dummy[predictors])  
y = np.array(math_data_dummy[outcome])
```

**2.6 creating training and testing data**

```
[263]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,  
↳ shuffle=True, random_state=random_seed)
```

**2.7 standardizing the data (z-scoring)** to avoid wildly different values in the model

```
[264]: z = StandardScaler() # standard scaling object  
  
# fitting scaler object to training data then outputting it  
X_train = z.fit_transform(X_train)  
# fit the testing data on the distributions found in the training data  
X_test = z.transform(X_test)
```

## 2.3 3 modeling

### 3.1 building initial model

```
[265]: model = models.Sequential()  
model.add(layers.Dense(32, activation='relu', input_shape=(X_train.shape[1],)))  
model.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers to  
↳ help with overfitting  
model.add(layers.Dense(16, activation='relu'))  
model.add(layers.Dense(1)) # adding this layer because it is a regression  
↳ problem
```

#### 3.1.1 compiling model loss

```
[266]: model.compile(optimizer='rmsprop',  
                    loss='mse',  
                    metrics=['mae'])
```

### 3.1.2 fitting model on training data

```
[267]: model_output = model.fit(X_train, y_train, batch_size=20, epochs=20)
```

```
Epoch 1/20  
316/316 [=====] - 0s 1ms/step - loss: 119.6751 - mae:  
9.9906  
Epoch 2/20  
316/316 [=====] - 0s 72us/step - loss: 101.0312 - mae:  
9.1572  
Epoch 3/20  
316/316 [=====] - 0s 64us/step - loss: 86.5112 - mae:  
8.4750  
Epoch 4/20  
316/316 [=====] - 0s 68us/step - loss: 72.8146 - mae:  
7.6959  
Epoch 5/20  
316/316 [=====] - 0s 69us/step - loss: 58.0657 - mae:  
6.8205  
Epoch 6/20  
316/316 [=====] - 0s 64us/step - loss: 47.2777 - mae:  
6.0455  
Epoch 7/20  
316/316 [=====] - 0s 64us/step - loss: 35.8962 - mae:  
5.1757  
Epoch 8/20  
316/316 [=====] - 0s 65us/step - loss: 28.4997 - mae:  
4.3849  
Epoch 9/20  
316/316 [=====] - 0s 60us/step - loss: 25.6380 - mae:  
4.1440  
Epoch 10/20  
316/316 [=====] - 0s 63us/step - loss: 25.0365 - mae:  
3.9596  
Epoch 11/20  
316/316 [=====] - 0s 62us/step - loss: 23.5457 - mae:  
3.8782  
Epoch 12/20  
316/316 [=====] - 0s 69us/step - loss: 21.7583 - mae:  
3.7079  
Epoch 13/20  
316/316 [=====] - 0s 63us/step - loss: 21.3804 - mae:  
3.5563  
Epoch 14/20
```

```

316/316 [=====] - 0s 63us/step - loss: 20.8165 - mae:
3.6028
Epoch 15/20
316/316 [=====] - 0s 64us/step - loss: 20.0247 - mae:
3.5211
Epoch 16/20
316/316 [=====] - 0s 66us/step - loss: 19.4162 - mae:
3.5992
Epoch 17/20
316/316 [=====] - 0s 62us/step - loss: 19.7602 - mae:
3.4783
Epoch 18/20
316/316 [=====] - 0s 61us/step - loss: 18.5260 - mae:
3.4555
Epoch 19/20
316/316 [=====] - 0s 65us/step - loss: 19.4620 - mae:
3.4493
Epoch 20/20
316/316 [=====] - 0s 65us/step - loss: 18.6162 - mae:
3.3783

```

### 3.1.3 predicting with test data

```

[268]: predicted = model.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted[0]}")
print(f"test: mean absolute error loss {predicted[1]}")

```

```

79/79 [=====] - 0s 934us/step
test: mean squared error loss 15.240999076939836
test: mean absolute error loss 3.09421443939209

```

**3.2 further experiments** The following experiments will help convince you that the architecture choices you've made are all fairly reasonable, although there's still room for improvement: 1. You used two hidden layers. Try using one or three hidden layers, and see how doing so affects validation and test accuracy. 2. Try using layers with more hidden units or fewer hidden units: 32 units, 64 units, and so on. 3. Try using the mae loss function instead of mse. 4. Try using the tanh activation (an activation that was popular in the early days of neural networks) instead of relu.

#### 3.2.1 model using one hidden layer

```

[269]: model_slim = models.Sequential()
model_slim.add(layers.Dense(32, activation='relu', input_shape=(X_train.
    ↳shape[1],)))
model_slim.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers_
    ↳to help with overfitting
model_slim.add(layers.Dense(1)) # adding this layer because it is a regression_
    ↳problem

```



```

model_slim.compile(optimizer='rmsprop',
                   loss='mse',
                   metrics=['mae'])

model_slim_output = model_slim.fit(X_train, y_train, batch_size=20, epochs=20)

```

```

Epoch 1/20
316/316 [=====] - 0s 876us/step - loss: 114.9323 - mae:
9.7166
Epoch 2/20
316/316 [=====] - 0s 60us/step - loss: 106.8649 - mae:
9.4300
Epoch 3/20
316/316 [=====] - 0s 59us/step - loss: 96.7718 - mae:
8.9687
Epoch 4/20
316/316 [=====] - 0s 55us/step - loss: 87.5347 - mae:
8.4870
Epoch 5/20
316/316 [=====] - 0s 86us/step - loss: 81.4040 - mae:
8.1444
Epoch 6/20
316/316 [=====] - 0s 77us/step - loss: 74.4145 - mae:
7.7717
Epoch 7/20
316/316 [=====] - 0s 59us/step - loss: 67.0496 - mae:
7.3437
Epoch 8/20
316/316 [=====] - 0s 60us/step - loss: 60.9565 - mae:
6.9774
Epoch 9/20
316/316 [=====] - 0s 54us/step - loss: 54.7326 - mae:
6.5678
Epoch 10/20
316/316 [=====] - 0s 58us/step - loss: 47.6141 - mae:
6.0871
Epoch 11/20
316/316 [=====] - 0s 55us/step - loss: 45.6499 - mae:
5.8310
Epoch 12/20
316/316 [=====] - 0s 56us/step - loss: 42.7735 - mae:
5.6728
Epoch 13/20
316/316 [=====] - 0s 57us/step - loss: 35.2782 - mae:
5.0216
Epoch 14/20
316/316 [=====] - 0s 58us/step - loss: 33.4923 - mae:
4.8642

```

```

Epoch 15/20
316/316 [=====] - 0s 57us/step - loss: 28.8793 - mae:
4.4988
Epoch 16/20
316/316 [=====] - 0s 56us/step - loss: 28.3410 - mae:
4.4444
Epoch 17/20
316/316 [=====] - 0s 60us/step - loss: 26.5100 - mae:
4.3034
Epoch 18/20
316/316 [=====] - 0s 57us/step - loss: 24.3877 - mae:
4.0470
Epoch 19/20
316/316 [=====] - 0s 56us/step - loss: 23.3842 - mae:
3.9314
Epoch 20/20
316/316 [=====] - 0s 56us/step - loss: 23.1853 - mae:
3.9243

```

```

[270]: predicted_slim = model_slim.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted_slim[0]}")
print(f"test: mean absolute error loss {predicted_slim[1]}")

```

```

79/79 [=====] - 0s 851us/step
test: mean squared error loss 22.277899392043487
test: mean absolute error loss 3.940617799758911

```

underfitting here, and not capturing the real values in the data.

### 3.2.2 model using more hidden units

```

[271]: model_fat = models.Sequential()
model_fat.add(layers.Dense(256, activation='relu', input_shape=(X_train.
↪shape[1],)))
model_fat.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers to ↵
↪help with overfitting
model_fat.add(layers.Dense(128, activation='relu'))
model_fat.add(layers.Dropout(0.3, seed=random_seed))
model_fat.add(layers.Dense(1)) # adding this layer because it is a regression ↵
↪problem

model_fat.compile(optimizer='rmsprop',
                  loss='mse',
                  metrics=['mae'])

model_fat_output = model_fat.fit(X_train, y_train, batch_size=20, epochs=20)

```

```

Epoch 1/20
316/316 [=====] - 0s 1ms/step - loss: 57.5891 - mae:

```

6.4724  
Epoch 2/20  
316/316 [=====] - 0s 244us/step - loss: 22.2892 - mae: 3.7171  
Epoch 3/20  
316/316 [=====] - 0s 251us/step - loss: 17.6542 - mae: 3.2800  
Epoch 4/20  
316/316 [=====] - 0s 260us/step - loss: 17.3445 - mae: 3.3105  
Epoch 5/20  
316/316 [=====] - 0s 264us/step - loss: 16.1278 - mae: 3.1216  
Epoch 6/20  
316/316 [=====] - 0s 261us/step - loss: 15.1920 - mae: 3.0420  
Epoch 7/20  
316/316 [=====] - 0s 245us/step - loss: 14.4454 - mae: 3.0478  
Epoch 8/20  
316/316 [=====] - 0s 234us/step - loss: 12.9277 - mae: 2.7721  
Epoch 9/20  
316/316 [=====] - 0s 246us/step - loss: 12.9546 - mae: 2.7830  
Epoch 10/20  
316/316 [=====] - 0s 241us/step - loss: 11.2297 - mae: 2.6013  
Epoch 11/20  
316/316 [=====] - 0s 239us/step - loss: 12.4494 - mae: 2.7967  
Epoch 12/20  
316/316 [=====] - 0s 239us/step - loss: 11.2532 - mae: 2.5223  
Epoch 13/20  
316/316 [=====] - 0s 234us/step - loss: 9.3080 - mae: 2.4208  
Epoch 14/20  
316/316 [=====] - 0s 236us/step - loss: 10.1783 - mae: 2.4969  
Epoch 15/20  
316/316 [=====] - 0s 245us/step - loss: 10.4879 - mae: 2.5408  
Epoch 16/20  
316/316 [=====] - 0s 235us/step - loss: 9.5909 - mae: 2.3872  
Epoch 17/20  
316/316 [=====] - 0s 274us/step - loss: 9.8122 - mae:

```

2.3864
Epoch 18/20
316/316 [=====] - 0s 247us/step - loss: 9.6088 - mae:
2.4541
Epoch 19/20
316/316 [=====] - 0s 257us/step - loss: 8.8385 - mae:
2.3844
Epoch 20/20
316/316 [=====] - 0s 259us/step - loss: 7.9328 - mae:
2.2292

```

```

[272]: predicted_fat = model_fat.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted_fat[0]}")
print(f"test: mean absolute error loss {predicted_fat[1]}")

```

```

79/79 [=====] - 0s 1ms/step
test: mean squared error loss 15.176022903828681
test: mean absolute error loss 2.9776809215545654

```

model performs good here, better than the original geometry

### 3.2.3 model using mae loss

```

[273]: model_mae = models.Sequential()
model_mae.add(layers.Dense(32, activation='relu', input_shape=(X_train.
    ↪shape[1],)))
model_mae.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers to
    ↪help with overfitting
model_mae.add(layers.Dense(16, activation='relu'))
model_mae.add(layers.Dense(1)) # adding this layer because it is a regression
    ↪problem

model_mae.compile(optimizer='rmsprop',
                  loss='mae',
                  metrics=['mse'])

model_mae_output = model_mae.fit(X_train, y_train, batch_size=20, epochs=20)

```

```

Epoch 1/20
316/316 [=====] - 0s 1ms/step - loss: 10.0771 - mse:
122.1471
Epoch 2/20
316/316 [=====] - 0s 72us/step - loss: 9.1377 - mse:
101.8730
Epoch 3/20
316/316 [=====] - 0s 64us/step - loss: 8.2996 - mse:
85.0169
Epoch 4/20
316/316 [=====] - 0s 72us/step - loss: 7.4872 - mse:

```

```

69.8086
Epoch 5/20
316/316 [=====] - 0s 66us/step - loss: 6.3562 - mse:
53.3596
Epoch 6/20
316/316 [=====] - 0s 67us/step - loss: 5.5196 - mse:
41.4658
Epoch 7/20
316/316 [=====] - 0s 67us/step - loss: 5.0455 - mse:
34.3769
Epoch 8/20
316/316 [=====] - 0s 61us/step - loss: 4.5097 - mse:
29.7655
Epoch 9/20
316/316 [=====] - 0s 64us/step - loss: 4.2765 - mse:
27.2567
Epoch 10/20
316/316 [=====] - 0s 61us/step - loss: 3.9621 - mse:
24.7270
Epoch 11/20
316/316 [=====] - 0s 66us/step - loss: 3.8453 - mse:
24.1783
Epoch 12/20
316/316 [=====] - 0s 62us/step - loss: 3.9050 - mse:
23.8684
Epoch 13/20
316/316 [=====] - 0s 64us/step - loss: 3.6350 - mse:
21.7358
Epoch 14/20
316/316 [=====] - 0s 63us/step - loss: 3.8395 - mse:
22.9980
Epoch 15/20
316/316 [=====] - 0s 67us/step - loss: 3.5856 - mse:
21.3374
Epoch 16/20
316/316 [=====] - 0s 62us/step - loss: 3.6338 - mse:
20.5292
Epoch 17/20
316/316 [=====] - 0s 63us/step - loss: 3.5319 - mse:
20.0993
Epoch 18/20
316/316 [=====] - 0s 66us/step - loss: 3.4101 - mse:
18.6605
Epoch 19/20
316/316 [=====] - 0s 64us/step - loss: 3.4450 - mse:
20.0430
Epoch 20/20
316/316 [=====] - 0s 62us/step - loss: 3.4815 - mse:

```

20.1830

```
[274]: predicted_mae_model = model_mae.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted_mae_model[1]}")
print(f"test: mean absolute error loss {predicted_mae_model[0]}")
```

```
79/79 [=====] - 0s 933us/step
test: mean squared error loss 16.141311645507812
test: mean absolute error loss 3.067770185349863
```

the model performs about the same in this case with a little bit better scores overall

### 3.2.4 model using tanh activation

```
[275]: model_tanh = models.Sequential()
model_tanh.add(layers.Dense(32, activation='tanh', input_shape=(X_train.
    ↪shape[1],)))
model_tanh.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers
    ↪to help with overfitting
model_tanh.add(layers.Dense(16, activation='tanh'))
model_tanh.add(layers.Dense(1)) # adding this layer because it is a regression
    ↪problem

model_tanh.compile(optimizer='rmsprop',
                    loss='mse',
                    metrics=['mae'])

model_tanh_output = model_tanh.fit(X_train, y_train, batch_size=20, epochs=20)
```

```
Epoch 1/20
316/316 [=====] - 0s 1ms/step - loss: 126.9281 - mae:
10.2410
Epoch 2/20
316/316 [=====] - 0s 70us/step - loss: 122.7916 - mae:
10.0862
Epoch 3/20
316/316 [=====] - 0s 66us/step - loss: 120.1034 - mae:
9.9677
Epoch 4/20
316/316 [=====] - 0s 64us/step - loss: 115.7618 - mae:
9.7880
Epoch 5/20
316/316 [=====] - 0s 70us/step - loss: 112.6738 - mae:
9.6679
Epoch 6/20
316/316 [=====] - 0s 64us/step - loss: 107.7759 - mae:
9.4378
Epoch 7/20
316/316 [=====] - 0s 65us/step - loss: 103.9476 - mae:
```

```

9.2530
Epoch 8/20
316/316 [=====] - 0s 64us/step - loss: 97.5534 - mae:
8.9889
Epoch 9/20
316/316 [=====] - 0s 65us/step - loss: 93.6230 - mae:
8.7832
Epoch 10/20
316/316 [=====] - 0s 63us/step - loss: 88.3186 - mae:
8.4938
Epoch 11/20
316/316 [=====] - 0s 63us/step - loss: 82.2173 - mae:
8.1885
Epoch 12/20
316/316 [=====] - 0s 63us/step - loss: 77.5650 - mae:
7.9221
Epoch 13/20
316/316 [=====] - 0s 67us/step - loss: 71.0205 - mae:
7.5801
Epoch 14/20
316/316 [=====] - 0s 62us/step - loss: 67.1885 - mae:
7.2732
Epoch 15/20
316/316 [=====] - 0s 64us/step - loss: 59.8460 - mae:
6.8081
Epoch 16/20
316/316 [=====] - 0s 67us/step - loss: 54.1449 - mae:
6.4695
Epoch 17/20
316/316 [=====] - 0s 64us/step - loss: 49.7382 - mae:
6.1439
Epoch 18/20
316/316 [=====] - 0s 65us/step - loss: 44.5213 - mae:
5.7813
Epoch 19/20
316/316 [=====] - 0s 64us/step - loss: 41.7015 - mae:
5.5089
Epoch 20/20
316/316 [=====] - 0s 65us/step - loss: 36.0837 - mae:
5.1028

```

```

[276]: predicted_tanh = model_tanh.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted_tanh[0]}")
print(f"test: mean absolute error loss {predicted_tanh[1]}")

```

```

79/79 [=====] - 0s 923us/step
test: mean squared error loss 44.83786131460455
test: mean absolute error loss 5.621185779571533

```

lots more loss in a model that uses a tanh activation vs relu activation for this neural network geometry

**3.3 k fold validation** the code here is taken out of the textbook for the regression examples

### 3.3.1 build model function definition

```
[277]: def build_model():
        model = models.Sequential()
        model.add(layers.Dense(32, activation='relu',
                                input_shape=(X_train.shape[1],)))
        model.add(layers.Dense(16, activation='relu'))
        model.add(layers.Dense(1))
        model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
        return model
```

### 3.3.2 defining k-fold variables

```
[278]: n_folds = 10
        num_epochs = 20
        num_val_samples = len(X_train) // n_folds
        all_mae_histories = []
```

### 3.3.3 k-fold loop

```
[279]: for i in range(n_folds):
        print("processing fold #", i)
        val_data = X_train[i * num_val_samples : (i + 1) * num_val_samples]
        val_targets = y_train[i * num_val_samples : (i + 1) * num_val_samples]

        partial_train_data = np.concatenate(
            [X_train[:i * num_val_samples],
             X_train[(i + 1) * num_val_samples:]],
            axis = 0)

        partial_train_targets = np.concatenate(
            [y_train[:i * num_val_samples],
             y_train[(i + 1) * num_val_samples:]],
            axis = 0)

        model = build_model()
        history = model.fit(partial_train_data, partial_train_targets,
                            validation_data=(val_data, val_targets),
                            epochs=num_epochs, batch_size=1, verbose=0)

        mae_history = history.history['val_mae']
        all_mae_histories.append(mae_history)
```




```
processing fold # 0
processing fold # 1
processing fold # 2
processing fold # 3
processing fold # 4
processing fold # 5
processing fold # 6
processing fold # 7
processing fold # 8
processing fold # 9
```

### 3.3.4 plotting k-fold results

```
[280]: average_mae_history = [
        np.mean([x[i] for x in all_mae_histories]) for i in range(num_epochs)]
```

```
[281]: import matplotlib.pyplot as plt
plt.plot(range(1, len(average_mae_history) + 1), average_mae_history)
plt.xlabel('Epochs')
plt.ylabel('Validation MAE')
plt.show()
```




looks like validation mean absolute error is minimized at around 10 epochs or so

```
[282]: def smooth_curve(points, factor=0.9):
        smoothed_points = []
```

```

for point in points:
    if smoothed_points:
        previous = smoothed_points[-1]
        smoothed_points.append(previous * factor + point * (1 - factor))
    else:
        smoothed_points.append(point)
return smoothed_points
smooth_mae_history = smooth_curve(average_mae_history[10:])
plt.plot(range(1, len(smooth_mae_history) + 1), smooth_mae_history)
plt.xlabel('Epochs')
plt.ylabel('Validation MAE')
plt.show()

```



this shows the average validation score vs number of epochs. it goes up over the epochs, but is not that different in absolute terms.

### 3.3.5 building model from validation

```

[283]: model = build_model()
model.fit(X_train, y_train,
          epochs = 10, batch_size = 20)

test_mse_score, test_mae_score = model.evaluate(X_test, y_test)
print(f"testing mean absolute error {test_mae_score}")

```

Epoch 1/10

```

316/316 [=====] - 0s 852us/step - loss: 123.0432 - mae:
10.0735
Epoch 2/10
316/316 [=====] - 0s 62us/step - loss: 109.4362 - mae:
9.5152
Epoch 3/10
316/316 [=====] - 0s 53us/step - loss: 96.4149 - mae:
8.9209
Epoch 4/10
316/316 [=====] - 0s 58us/step - loss: 82.2523 - mae:
8.2291
Epoch 5/10
316/316 [=====] - 0s 59us/step - loss: 67.8744 - mae:
7.4199
Epoch 6/10
316/316 [=====] - 0s 57us/step - loss: 53.8478 - mae:
6.5491
Epoch 7/10
316/316 [=====] - 0s 54us/step - loss: 41.5124 - mae:
5.6553
Epoch 8/10
316/316 [=====] - 0s 55us/step - loss: 31.3870 - mae:
4.8664
Epoch 9/10
316/316 [=====] - 0s 54us/step - loss: 24.1829 - mae:
4.1737
Epoch 10/10
316/316 [=====] - 0s 52us/step - loss: 20.1880 - mae:
3.7080
79/79 [=====] - 0s 666us/step
testing mean absolute error 3.905726909637451

```

the mean absolute error for the testing error is 3.3 grade units out of 20. this model does pretty well and is not far off from the training data. 3.26 mae in training and 3.3 in testing. epochs of 10 are a good way to normalize it for this network geometry.