

# CLydonAssignment2

March 2, 2022

## 1 Assignment 2

- 1) Read chapters 2-4 in your textbook and review the coding examples we went over in class
- 2) Review the Keras documentation for things like the Layer types and Optimizer types to better familiarize yourself
- 3) Redo the coding examples and do the “Further experiments” in chapter 3 of your book with datasets you find interesting and may want to use for further assignments and projects
- 4) Provide a brief write up of the experiment and analysis.

## 2 table of contents

1. import libraries
2. data prep
  1. reading data
  2. finding summary data
  3. dummyizing data
  4. creating predictor & outcome attribute lists
  5. transforming data into numpy arrays
  6. creating training and testing data
  7. standardizing the data (z-scoring)
3. modeling
  1. initial model
    1. building layers
    2. fitting model on training data
    3. predicting with test data
  2. further experiments
    1. model using one hiddden layer
    2. model using more hidden units
    3. model using mae loss
    4. model using tanh activation
  3. kfold optimization
    1. build model function definition
    2. defining k-fold variables
    3. k-fold loop
    4. plotting k-fold results
    5. building model from validation

---

## 2.1 1 import libraries

```
[1]: # data prep
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.model_selection import train_test_split

import numpy as np
import pandas as pd

# keras imports
from keras import models
from keras import layers
from keras import optimizers

# random seed for pseudo-random operations
random_seed = 100
```

Using TensorFlow backend.

---

## 2.2 2 data prep

### 2.1 reading dataset

```
[2]: math_data = pd.read_csv("https://raw.githubusercontent.com/connoralydon/
↳datasets/main/student-math.csv", sep=";")
```

data can be found [here](#)

Attributes for both student-mat.csv (Math course) and student-por.csv (Portuguese language course) datasets: 1 school - student's school (binary: 'GP' - Gabriel Pereira or 'MS' - Mousinho da Silveira) 2 sex - student's sex (binary: 'F' - female or 'M' - male) 3 age - student's age (numeric: from 15 to 22) 4 address - student's home address type (binary: 'U' - urban or 'R' - rural) 5 famsize - family size (binary: 'LE3' - less or equal to 3 or 'GT3' - greater than 3) 6 Pstatus - parent's cohabitation status (binary: 'T' - living together or 'A' - apart) 7 Medu - mother's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education) 8 Fedu - father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education) 9 Mjob - mother's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at\_home' or 'other') 10 Fjob - father's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at\_home' or 'other') 11 reason - reason to choose this school (nominal: close to 'home', school 'reputation', 'course' preference or 'other') 12 guardian - student's guardian (nominal: 'mother', 'father' or 'other') 13 traveltime - home to school travel time (numeric: 1 - <15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >1 hour) 14 studytime - weekly study time (numeric: 1 - <2 hours, 2 - 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours) 15 failures - number of past class failures (numeric: n if 1 ≤ n < 3, else 4) 16 schoolsup - extra educational support (binary: yes or no) 17 famsup - family educational support (binary: yes or no) 18 paid - extra paid classes within the course subject (Math or Portuguese) (binary: yes

or no) 19 activities - extra-curricular activities (binary: yes or no) 20 nursery - attended nursery school (binary: yes or no) 21 higher - wants to take higher education (binary: yes or no) 22 internet - Internet access at home (binary: yes or no) 23 romantic - with a romantic relationship (binary: yes or no) 24 famrel - quality of family relationships (numeric: from 1 - very bad to 5 - excellent) 25 freetime - free time after school (numeric: from 1 - very low to 5 - very high) 26 goout - going out with friends (numeric: from 1 - very low to 5 - very high) 27 Dalc - workday alcohol consumption (numeric: from 1 - very low to 5 - very high) 28 Walc - weekend alcohol consumption (numeric: from 1 - very low to 5 - very high) 29 health - current health status (numeric: from 1 - very bad to 5 - very good) 30 absences - number of school absences (numeric: from 0 to 93) these grades are related with the course subject, Math or Portuguese: 31 G1 - first period grade (numeric: from 0 to 20) 32 G2 - second period grade (numeric: from 0 to 20) 33 G3 - final grade (numeric: from 0 to 20, output target)

## 2.2 summary info about data

```
[3]: math_data.head
```

```
[3]: <bound method NDFrame.head of
Fedu      Mjob      Fjob      \
0         GP      F      18      U      GT3      A      4      4      at_home      teacher
1         GP      F      17      U      GT3      T      1      1      at_home      other
2         GP      F      15      U      LE3      T      1      1      at_home      other
3         GP      F      15      U      GT3      T      4      2      health      services
4         GP      F      16      U      GT3      T      3      3      other      other
...      ...      ...      ...      ...      ...      ...      ...      ...      ...
390      MS      M      20      U      LE3      A      2      2      services      services
391      MS      M      17      U      LE3      T      3      1      services      services
392      MS      M      21      R      GT3      T      1      1      other      other
393      MS      M      18      R      LE3      T      3      2      services      other
394      MS      M      19      U      LE3      T      1      1      other      at_home

      ...      famrel      freetime      goout      Dalc      Walc      health      absences      G1      G2      G3
0      ...      4      3      4      1      1      3      6      5      6      6
1      ...      5      3      3      1      1      3      4      5      5      6
2      ...      4      3      2      2      3      3      10      7      8      10
3      ...      3      2      2      1      1      5      2      15      14      15
4      ...      4      3      2      1      2      5      4      6      10      10
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
390      ...      5      5      4      4      5      4      11      9      9      9
391      ...      2      4      5      3      4      2      3      14      16      16
392      ...      5      5      3      3      3      3      3      10      8      7
393      ...      4      4      1      3      4      5      0      11      12      10
394      ...      3      2      3      3      3      5      5      8      9      9
```

```
[395 rows x 33 columns]>
```

```
[4]: math_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 395 entries, 0 to 394

Data columns (total 33 columns):

#	Column	Non-Null Count	Dtype
0	school	395 non-null	object
1	sex	395 non-null	object
2	age	395 non-null	int64
3	address	395 non-null	object
4	famsize	395 non-null	object
5	Pstatus	395 non-null	object
6	Medu	395 non-null	int64
7	Fedu	395 non-null	int64
8	Mjob	395 non-null	object
9	Fjob	395 non-null	object
10	reason	395 non-null	object
11	guardian	395 non-null	object
12	traveltime	395 non-null	int64
13	studytime	395 non-null	int64
14	failures	395 non-null	int64
15	schoolsup	395 non-null	object
16	famsup	395 non-null	object
17	paid	395 non-null	object
18	activities	395 non-null	object
19	nursery	395 non-null	object
20	higher	395 non-null	object
21	internet	395 non-null	object
22	romantic	395 non-null	object
23	famrel	395 non-null	int64
24	freetime	395 non-null	int64
25	goout	395 non-null	int64
26	Dalc	395 non-null	int64
27	Walc	395 non-null	int64
28	health	395 non-null	int64
29	absences	395 non-null	int64
30	G1	395 non-null	int64
31	G2	395 non-null	int64
32	G3	395 non-null	int64

dtypes: int64(16), object(17)

memory usage: 102.0+ KB

```
[5]: math_data.dtypes
```

```
[5]: school      object
     sex         object
     age         int64
     address     object
     famsize     object
```

```

Pstatus      object
Medu         int64
Fedu         int64
Mjob         object
Fjob         object
reason       object
guardian     object
traveltime   int64
studytime    int64
failures     int64
schoolsup    object
famsup       object
paid         object
activities   object
nursery      object
higher       object
internet     object
romantic     object
famrel       int64
freetime     int64
goout        int64
Dalc         int64
Walc         int64
health       int64
absences     int64
G1           int64
G2           int64
G3           int64
dtype: object

```

```
[6]: math_data.columns
```

```

[6]: Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
          'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
          'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
          'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
          'Walc', 'health', 'absences', 'G1', 'G2', 'G3'],
          dtype='object')

```

```
[7]: math_data.shape
```

```
[7]: (395, 33)
```

```

[8]: # observe the different values used for species - 2 different values: 0, 1
def print_range(data):
    print(f"range of {data.min()} to {data.max()}\n")

```

```
print("First test")
print_range(math_data["G2"])

print("Second test")
print_range(math_data["G1"])

print("Final test")
print_range(math_data["G3"])
```

First test  
range of 0 to 19

Second test  
range of 3 to 19

Final test  
range of 0 to 20

```
[9]: print("total null values")
math_data.isnull().sum(axis=0).sum()
```

total null values

[9]: 0

```
[10]: print("num duplicates")
math_data.duplicated().sum(axis=0)
```

num duplicates

[10]: 0

### 2.3 dummyizing data

```
[11]: categorical_cols = {}

for col in math_data.columns:
    if(math_data[col].dtype == "object"):
        categorical_cols[col] = list(math_data[col].unique())

numerical_cols = list(set(math_data.columns) - set(categorical_cols.keys())) #_
↳ set difference

math_data_dummy = pd.get_dummies(math_data, columns = categorical_cols.keys())
```

**2.4 model variables** removing g1, g2, g3 because they would be highly autocorrelary. someone who does good on one test is likely to do good on another.

```
[12]: predictors = list(math_data_dummy.columns)
predictors.remove("G1") # first exam
predictors.remove("G2") # second exam
predictors.remove("G3") # third exam

outcome = "G3"
```

**2.5 transforming data into numpy arrays** this is because keras won't take pandas dataframes

```
[13]: X = np.array(math_data_dummy[predictors])
y = np.array(math_data_dummy[outcome])
```

**2.6 creating training and testing data**

```
[14]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,
↳shuffle=True, random_state=random_seed)
```

**2.7 standardizing the data (z-scoring)** to avoid wildly different values in the model

```
[15]: z = StandardScaler() # standard scaling object
# z = MinMaxScaler() # alternative min max scaler so all values should lie
↳between 0-1 for training and about that for testing

# fitting scaler object to training data then outputting it
X_train = z.fit_transform(X_train)
# fit the testing data on the distributions found in the training data
X_test = z.transform(X_test)
```

## 2.3 3 modeling

### 3.1 building initial model

```
[16]: model = models.Sequential()
model.add(layers.Dense(32, activation='relu', input_shape=(X_train.shape[1],)))
model.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers to
↳help with overfitting
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1)) # adding this layer because it is a regression
↳problem
```

```
2022-03-02 00:32:31.494001: I tensorflow/core/platform/cpu_feature_guard.cc:145]
This TensorFlow binary is optimized with Intel(R) MKL-DNN to use the following
CPU instructions in performance critical operations: SSE4.1 SSE4.2
To enable them in non-MKL-DNN operations, rebuild TensorFlow with the
appropriate compiler flags.
2022-03-02 00:32:31.495570: I
tensorflow/core/common_runtime/process_util.cc:115] Creating new thread pool
```

with default inter op setting: 8. Tune using `inter_op_parallelism_threads` for best performance.

### *3.1.1 compiling model loss*

```
[17]: model.compile(optimizer='rmsprop',  
                  loss='mse',  
                  metrics=['mae'])
```

### *3.1.2 fitting model on training data*

```
[18]: model_output = model.fit(X_train, y_train, batch_size=20, epochs=20)
```

```
Epoch 1/20  
316/316 [=====] - 0s 1ms/step - loss: 118.3580 - mae:  
9.9130  
Epoch 2/20  
316/316 [=====] - 0s 64us/step - loss: 102.6289 - mae:  
9.2095  
Epoch 3/20  
316/316 [=====] - 0s 69us/step - loss: 91.5656 - mae:  
8.6728  
Epoch 4/20  
316/316 [=====] - 0s 68us/step - loss: 76.0432 - mae:  
7.9212  
Epoch 5/20  
316/316 [=====] - 0s 84us/step - loss: 63.1482 - mae:  
7.1596  
Epoch 6/20  
316/316 [=====] - 0s 64us/step - loss: 51.6425 - mae:  
6.3279  
Epoch 7/20  
316/316 [=====] - 0s 64us/step - loss: 39.1787 - mae:  
5.4857  
Epoch 8/20  
316/316 [=====] - 0s 65us/step - loss: 29.8894 - mae:  
4.6153  
Epoch 9/20  
316/316 [=====] - 0s 62us/step - loss: 22.9944 - mae:  
3.9381  
Epoch 10/20  
316/316 [=====] - 0s 64us/step - loss: 24.5114 - mae:  
3.9499  
Epoch 11/20  
316/316 [=====] - 0s 67us/step - loss: 20.8206 - mae:  
3.6482  
Epoch 12/20  
316/316 [=====] - 0s 71us/step - loss: 20.5493 - mae:  
3.5286  
Epoch 13/20
```



```

316/316 [=====] - 0s 73us/step - loss: 18.8742 - mae:
3.5030
Epoch 14/20
316/316 [=====] - 0s 70us/step - loss: 19.9681 - mae:
3.5331
Epoch 15/20
316/316 [=====] - 0s 68us/step - loss: 18.8069 - mae:
3.4624
Epoch 16/20
316/316 [=====] - 0s 94us/step - loss: 18.0441 - mae:
3.3671
Epoch 17/20
316/316 [=====] - 0s 108us/step - loss: 19.3649 - mae:
3.4647
Epoch 18/20
316/316 [=====] - 0s 90us/step - loss: 18.2037 - mae:
3.2381
Epoch 19/20
316/316 [=====] - 0s 85us/step - loss: 18.5090 - mae:
3.3865
Epoch 20/20
316/316 [=====] - 0s 71us/step - loss: 18.5258 - mae:
3.4294

```

### 3.1.3 predicting with test data

```

[19]: predicted = model.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted[0]}")
print(f"test: mean absolute error loss {predicted[1]}")

```

```

79/79 [=====] - 0s 973us/step
test: mean squared error loss 15.218327860288982
test: mean absolute error loss 3.1643083095550537

```

**3.2 further experiments** The following experiments will help convince you that the architecture choices you've made are all fairly reasonable, although there's still room for improvement: 1. You used two hidden layers. Try using one or three hidden layers, and see how doing so affects validation and test accuracy. 2. Try using layers with more hidden units or fewer hidden units: 32 units, 64 units, and so on. 3. Try using the mae loss function instead of mse. 4. Try using the tanh activation (an activation that was popular in the early days of neural networks) instead of relu.

#### 3.2.1 model using one hidden layer

```

[20]: model_slim = models.Sequential()
model_slim.add(layers.Dense(32, activation='relu', input_shape=(X_train.
    ↳ shape[1],)))
model_slim.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers_
    ↳ to help with overfitting

```

```

model_slim.add(layers.Dense(1)) # adding this layer because it is a regression
                                ↪problem

model_slim.compile(optimizer='rmsprop',
                   loss='mse',
                   metrics=['mae'])

model_slim_output = model_slim.fit(X_train, y_train, batch_size=20, epochs=20)

```

```

Epoch 1/20
316/316 [=====] - 0s 928us/step - loss: 130.8246 - mae:
10.4379
Epoch 2/20
316/316 [=====] - 0s 70us/step - loss: 120.3331 - mae:
9.9618
Epoch 3/20
316/316 [=====] - 0s 59us/step - loss: 113.6022 - mae:
9.6726
Epoch 4/20
316/316 [=====] - 0s 59us/step - loss: 104.0589 - mae:
9.2243
Epoch 5/20
316/316 [=====] - 0s 59us/step - loss: 96.9397 - mae:
8.8806
Epoch 6/20
316/316 [=====] - 0s 59us/step - loss: 89.3699 - mae:
8.5300
Epoch 7/20
316/316 [=====] - 0s 59us/step - loss: 82.7400 - mae:
8.1879
Epoch 8/20
316/316 [=====] - 0s 58us/step - loss: 73.7910 - mae:
7.6626
Epoch 9/20
316/316 [=====] - 0s 58us/step - loss: 66.5331 - mae:
7.3049
Epoch 10/20
316/316 [=====] - 0s 61us/step - loss: 61.9143 - mae:
7.0152
Epoch 11/20
316/316 [=====] - 0s 56us/step - loss: 55.4966 - mae:
6.5787
Epoch 12/20
316/316 [=====] - 0s 60us/step - loss: 50.5342 - mae:
6.2335
Epoch 13/20
316/316 [=====] - 0s 60us/step - loss: 45.1718 - mae:

```

```

5.8720
Epoch 14/20
316/316 [=====] - 0s 54us/step - loss: 43.4684 - mae:
5.7515
Epoch 15/20
316/316 [=====] - 0s 58us/step - loss: 39.0597 - mae:
5.3753
Epoch 16/20
316/316 [=====] - 0s 55us/step - loss: 31.9364 - mae:
4.7656
Epoch 17/20
316/316 [=====] - 0s 53us/step - loss: 30.7910 - mae:
4.6785
Epoch 18/20
316/316 [=====] - 0s 61us/step - loss: 27.2381 - mae:
4.2940
Epoch 19/20
316/316 [=====] - 0s 95us/step - loss: 24.6287 - mae:
4.1543
Epoch 20/20
316/316 [=====] - 0s 78us/step - loss: 23.8875 - mae:
3.9925

```

```

[21]: predicted_slim = model_slim.evaluate(X_test, y_test)
      print(f"test: mean squared error loss {predicted_slim[0]}")
      print(f"test: mean absolute error loss {predicted_slim[1]}")

```

```

79/79 [=====] - 0s 939us/step
test: mean squared error loss 25.2938937175123
test: mean absolute error loss 4.15211820602417

```

underfitting here, and not capturing the real values in the data.

### 3.2.2 model using more hidden units

```

[22]: model_fat = models.Sequential()
      model_fat.add(layers.Dense(256, activation='relu', input_shape=(X_train.
      ↪shape[1],)))
      model_fat.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers to
      ↪help with overfitting
      model_fat.add(layers.Dense(128, activation='relu'))
      model_fat.add(layers.Dropout(0.3, seed=random_seed))
      model_fat.add(layers.Dense(1)) # adding this layer because it is a regression
      ↪problem

      model_fat.compile(optimizer='rmsprop',
                        loss='mse',
                        metrics=['mae'])

```

```
model_fat_output = model_fat.fit(X_train, y_train, batch_size=20, epochs=20)
```

Epoch 1/20

316/316 [=====] - 0s 1ms/step - loss: 56.0610 - mae: 6.3596

Epoch 2/20

316/316 [=====] - 0s 243us/step - loss: 22.5144 - mae: 3.7731

Epoch 3/20

316/316 [=====] - 0s 281us/step - loss: 18.8599 - mae: 3.3756

Epoch 4/20

316/316 [=====] - 0s 319us/step - loss: 17.4967 - mae: 3.2330

Epoch 5/20

316/316 [=====] - 0s 256us/step - loss: 16.1123 - mae: 3.0867

Epoch 6/20

316/316 [=====] - 0s 277us/step - loss: 15.4312 - mae: 3.0866

Epoch 7/20

316/316 [=====] - 0s 252us/step - loss: 14.0872 - mae: 2.9112

Epoch 8/20

316/316 [=====] - 0s 243us/step - loss: 14.8829 - mae: 3.0200

Epoch 9/20

316/316 [=====] - 0s 270us/step - loss: 12.0349 - mae: 2.6493

Epoch 10/20

316/316 [=====] - 0s 287us/step - loss: 12.8396 - mae: 2.7926

Epoch 11/20

316/316 [=====] - 0s 294us/step - loss: 11.4187 - mae: 2.6157

Epoch 12/20

316/316 [=====] - 0s 260us/step - loss: 11.1469 - mae: 2.5922

Epoch 13/20

316/316 [=====] - 0s 260us/step - loss: 10.9802 - mae: 2.5253

Epoch 14/20

316/316 [=====] - 0s 301us/step - loss: 10.2873 - mae: 2.4899

Epoch 15/20

316/316 [=====] - 0s 350us/step - loss: 9.5828 - mae: 2.4347

```
Epoch 16/20
316/316 [=====] - 0s 233us/step - loss: 10.6177 - mae:
2.6469
Epoch 17/20
316/316 [=====] - 0s 258us/step - loss: 9.0591 - mae:
2.3782
Epoch 18/20
316/316 [=====] - 0s 234us/step - loss: 9.5922 - mae:
2.4089
Epoch 19/20
316/316 [=====] - 0s 240us/step - loss: 8.4024 - mae:
2.2935
Epoch 20/20
316/316 [=====] - 0s 251us/step - loss: 8.0401 - mae:
2.2003
```

```
[23]: predicted_fat = model_fat.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted_fat[0]}")
print(f"test: mean absolute error loss {predicted_fat[1]}")
```

```
79/79 [=====] - 0s 1ms/step
test: mean squared error loss 15.87204266801665
test: mean absolute error loss 3.047870635986328
```

model performs good here, better than the original geometry

### 3.2.3 model using mae loss

```
[24]: model_mae = models.Sequential()
model_mae.add(layers.Dense(32, activation='relu', input_shape=(X_train.
    ↪shape[1],)))
model_mae.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers to
    ↪help with overfitting
model_mae.add(layers.Dense(16, activation='relu'))
model_mae.add(layers.Dense(1)) # adding this layer because it is a regression
    ↪problem

model_mae.compile(optimizer='rmsprop',
                  loss='mae',
                  metrics=['mse'])

model_mae_output = model_mae.fit(X_train, y_train, batch_size=20, epochs=20)
```

```
Epoch 1/20
316/316 [=====] - 0s 1ms/step - loss: 9.6232 - mse:
112.5280
Epoch 2/20
316/316 [=====] - 0s 81us/step - loss: 8.7322 - mse:
92.5440
```

Epoch 3/20  
316/316 [=====] - 0s 93us/step - loss: 7.9136 - mse:  
77.4263

Epoch 4/20  
316/316 [=====] - 0s 84us/step - loss: 6.8690 - mse:  
60.8901

Epoch 5/20  
316/316 [=====] - 0s 102us/step - loss: 5.8822 - mse:  
46.2384

Epoch 6/20  
316/316 [=====] - 0s 107us/step - loss: 5.2873 - mse:  
38.4869

Epoch 7/20  
316/316 [=====] - 0s 62us/step - loss: 4.3899 - mse:  
28.6425

Epoch 8/20  
316/316 [=====] - 0s 66us/step - loss: 4.1893 - mse:  
25.4656

Epoch 9/20  
316/316 [=====] - 0s 61us/step - loss: 4.0343 - mse:  
24.7666

Epoch 10/20  
316/316 [=====] - 0s 65us/step - loss: 3.9112 - mse:  
24.5084

Epoch 11/20  
316/316 [=====] - 0s 62us/step - loss: 3.7142 - mse:  
23.0018

Epoch 12/20  
316/316 [=====] - 0s 69us/step - loss: 3.4378 - mse:  
20.2887

Epoch 13/20  
316/316 [=====] - 0s 61us/step - loss: 3.5241 - mse:  
20.3639

Epoch 14/20  
316/316 [=====] - 0s 62us/step - loss: 3.5881 - mse:  
21.1820

Epoch 15/20  
316/316 [=====] - 0s 65us/step - loss: 3.6313 - mse:  
21.2265

Epoch 16/20  
316/316 [=====] - 0s 64us/step - loss: 3.4816 - mse:  
19.9230

Epoch 17/20  
316/316 [=====] - 0s 62us/step - loss: 3.4035 - mse:  
20.3702

Epoch 18/20  
316/316 [=====] - 0s 62us/step - loss: 3.4672 - mse:  
20.6029

```
Epoch 19/20
316/316 [=====] - 0s 64us/step - loss: 3.4079 - mse:
19.7452
Epoch 20/20
316/316 [=====] - 0s 66us/step - loss: 3.4679 - mse:
20.1896
```

```
[25]: predicted_mae_model = model_mae.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted_mae_model[1]}")
print(f"test: mean absolute error loss {predicted_mae_model[0]}")
```

```
79/79 [=====] - 0s 944us/step
test: mean squared error loss 15.803775787353516
test: mean absolute error loss 3.0256048214586477
```

the model performs about the same in this case with a little bit better scores overall

### 3.2.4 model using tanh activation

```
[26]: model_tanh = models.Sequential()
model_tanh.add(layers.Dense(32, activation='tanh', input_shape=(X_train.
    ↳shape[1],)))
model_tanh.add(layers.Dropout(0.3, seed=random_seed)) # adding dropout layers_
    ↳to help with overfitting
model_tanh.add(layers.Dense(16, activation='tanh'))
model_tanh.add(layers.Dense(1)) # adding this layer because it is a regression_
    ↳problem

model_tanh.compile(optimizer='rmsprop',
                    loss='mse',
                    metrics=['mae'])

model_tanh_output = model_tanh.fit(X_train, y_train, batch_size=20, epochs=20)
```

```
Epoch 1/20
316/316 [=====] - 0s 1ms/step - loss: 126.5751 - mae:
10.2939
Epoch 2/20
316/316 [=====] - 0s 63us/step - loss: 122.1501 - mae:
10.0991
Epoch 3/20
316/316 [=====] - 0s 64us/step - loss: 117.3921 - mae:
9.9126
Epoch 4/20
316/316 [=====] - 0s 70us/step - loss: 114.7226 - mae:
9.8298
Epoch 5/20
316/316 [=====] - 0s 65us/step - loss: 111.6152 - mae:
9.6545
```

Epoch 6/20  
316/316 [=====] - 0s 63us/step - loss: 106.4805 - mae: 9.4421  
Epoch 7/20  
316/316 [=====] - 0s 59us/step - loss: 104.1886 - mae: 9.3610  
Epoch 8/20  
316/316 [=====] - 0s 62us/step - loss: 96.4676 - mae: 9.0145  
Epoch 9/20  
316/316 [=====] - 0s 58us/step - loss: 91.6401 - mae: 8.7799  
Epoch 10/20  
316/316 [=====] - 0s 77us/step - loss: 85.0526 - mae: 8.3990  
Epoch 11/20  
316/316 [=====] - 0s 70us/step - loss: 81.4942 - mae: 8.2401  
Epoch 12/20  
316/316 [=====] - 0s 58us/step - loss: 76.5033 - mae: 7.9346  
Epoch 13/20  
316/316 [=====] - 0s 57us/step - loss: 68.4408 - mae: 7.4645  
Epoch 14/20  
316/316 [=====] - 0s 60us/step - loss: 63.3529 - mae: 7.1666  
Epoch 15/20  
316/316 [=====] - 0s 60us/step - loss: 58.1826 - mae: 6.8303  
Epoch 16/20  
316/316 [=====] - 0s 78us/step - loss: 53.2556 - mae: 6.4459  
Epoch 17/20  
316/316 [=====] - 0s 92us/step - loss: 45.4849 - mae: 6.0320  
Epoch 18/20  
316/316 [=====] - 0s 57us/step - loss: 40.1577 - mae: 5.5182  
Epoch 19/20  
316/316 [=====] - 0s 59us/step - loss: 38.9120 - mae: 5.3465  
Epoch 20/20  
316/316 [=====] - 0s 61us/step - loss: 33.3191 - mae: 4.9218



```
[27]: predicted_tanh = model_tanh.evaluate(X_test, y_test)
print(f"test: mean squared error loss {predicted_tanh[0]}")
print(f"test: mean absolute error loss {predicted_tanh[1]}")
```

79/79 [=====] - 0s 926us/step

test: mean squared error loss 43.97593906257726

test: mean absolute error loss 5.622044086456299

lots more loss in a model that uses a tanh activation vs relu activation for this neural network geometry

**3.3 k fold validation** the code here is taken out of the textbook for the regression examples

### 3.3.1 build model function definition

```
[28]: def build_model():
    model = models.Sequential()
    model.add(layers.Dense(32, activation='relu',
                           input_shape=(X_train.shape[1],)))
    model.add(layers.Dense(16, activation='relu'))
    model.add(layers.Dense(1))
    model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
    return model
```

### 3.3.2 defining k-fold variables

```
[29]: n_folds = 10
num_epochs = 20
num_val_samples = len(X_train) // n_folds
all_mae_histories = []
```

### 3.3.3 k-fold loop

```
[30]: for i in range(n_folds):
    print("processing fold #", i)
    val_data = X_train[i * num_val_samples : (i + 1) * num_val_samples]
    val_targets = y_train[i * num_val_samples : (i + 1) * num_val_samples]

    partial_train_data = np.concatenate(
        [X_train[:i * num_val_samples],
         X_train[(i + 1) * num_val_samples:]],
        axis = 0)

    partial_train_targets = np.concatenate(
        [y_train[:i * num_val_samples],
         y_train[(i + 1) * num_val_samples:]],
        axis = 0)

    model = build_model()
    history = model.fit(partial_train_data, partial_train_targets,
```

```

validation_data=(val_data, val_targets),
epochs=num_epochs, batch_size=1, verbose=0)

mae_history = history.history['val_mae']
all_mae_histories.append(mae_history)

```

```

processing fold # 0
processing fold # 1
processing fold # 2
processing fold # 3
processing fold # 4
processing fold # 5
processing fold # 6
processing fold # 7
processing fold # 8
processing fold # 9

```

### 3.3.4 plotting k-fold results

```


[31]: average_mae_history = [
      np.mean([x[i] for x in all_mae_histories]) for i in range(num_epochs)]

```

```


[32]: import matplotlib.pyplot as plt
      plt.plot(range(1, len(average_mae_history) + 1), average_mae_history)
      plt.xlabel('Epochs')
      plt.ylabel('Validation MAE')
      plt.show()

```



looks like validation mean absolute error is minimized at around 10 epochs or so

```
[33]: def smooth_curve(points, factor=0.9):
    smoothed_points = []
    for point in points:
        if smoothed_points:
            previous = smoothed_points[-1]
            smoothed_points.append(previous * factor + point * (1 - factor))
        else:
            smoothed_points.append(point)
    return smoothed_points
smooth_mae_history = smooth_curve(average_mae_history[10:])
plt.plot(range(1, len(smooth_mae_history) + 1), smooth_mae_history)
plt.xlabel('Epochs')
plt.ylabel('Validation MAE')
plt.show()
```



this shows the average validation score vs number of epochs. it goes up over the epochs, but is not that different in absolute terms.

### *3.3.5 building model from validation*

```
[34]: model = build_model()
      model.fit(X_train, y_train,
                epochs = 10, batch_size = 20)

      test_mse_score, test_mae_score = model.evaluate(X_test, y_test)
      print(f"testing mean absolute error {test_mae_score}")
```

Epoch 1/10  
316/316 [=====] - 0s 899us/step - loss: 125.9590 - mae: 10.2354  
Epoch 2/10  
316/316 [=====] - 0s 62us/step - loss: 108.1365 - mae: 9.4563  
Epoch 3/10  
316/316 [=====] - 0s 77us/step - loss: 93.4494 - mae: 8.7767  
Epoch 4/10  
316/316 [=====] - 0s 54us/step - loss: 78.9494 - mae: 8.0436  
Epoch 5/10  
316/316 [=====] - 0s 67us/step - loss: 64.9423 - mae: 7.2385  
Epoch 6/10  
316/316 [=====] - 0s 60us/step - loss: 51.8186 - mae: 6.3765  
Epoch 7/10  
316/316 [=====] - 0s 58us/step - loss: 40.3951 - mae: 5.5143  
Epoch 8/10  
316/316 [=====] - 0s 61us/step - loss: 31.0415 - mae: 4.6974  
Epoch 9/10  
316/316 [=====] - 0s 55us/step - loss: 24.1974 - mae: 4.0408  
Epoch 10/10  
316/316 [=====] - 0s 74us/step - loss: 20.1712 - mae: 3.6308  
79/79 [=====] - 0s 723us/step  
testing mean absolute error 3.7550230026245117

the mean absolute error for the testing error is 3.3 grade units out of 20. this model does pretty well and is not far off from the training data. 3.26 mae in training and 3.3 in testing. epochs of 10 are a good way to normalize it for this network geometry.